

The Pyramid Principle

A working guide to structuring arguments for senior audiences

Barbara Minto's method · An applied guide with examples from an analytics function · August 2026

0. How to use this guide

This is not a summary of Barbara Minto's book. It is an independent guide to applying the method: it explains the ideas in its own words, with examples built around real situations in an analytics function inside a bank — risk committees, memos, model notes, escalations, difficult conversations — because the original material uses corporate examples from another era that make the transfer harder than it needs to be.

It is written to be read alone and also to be worked through as a team. Part I is the method. Part II applies it to eight recurring situations in any function that depends on others for its data and on a corporate centre for its decisions. Part III is practice, including a workshop agenda if you plan to teach it.

The method is not learned by reading. It is learned by rewriting. Section 24 is a one-page checklist to keep open while you write; section 25 is a 30-day practice plan. If you only use two parts of this guide, use those.

The rule that summarises the whole guide: if your reader has to reach the fourth paragraph to find out what you are asking for, you have already lost. The conclusion goes first; everything else exists to support it.

1. Contents

0. How to use this guide

Part I — The method

2. Why the pyramid: the underlying argument
3. Anatomy of the pyramid
4. MECE: the discipline of grouping
5. The introduction: the SCQA pattern
6. Building the pyramid: two routes
7. Ordering ideas within a group
8. Summarising groups: where most documents are lost
9. Deduction and induction
10. Before structuring: defining the problem properly
11. Getting the pyramid onto the page
12. The pyramid in presentations
13. Adapting to a banking and regulatory environment

Part II — Applied cases

14. Three worked examples
15. Data quality when someone else owns the data
16. Getting a decision in a globally centralised organisation
17. Building a business case for funding
18. Difficult conversation downward: asking someone on your team to change a behaviour

- 19. Difficult conversation upward: when the problematic behaviour is your manager's
- 20. Structuring a "no"
- 21. Handing over a project or a system
- 22. Explaining a technical subject in plain terms

Part III — Practice

- 23. Common failures in a quantitative profile
- 24. Checklist
- 25. A 30-day practice plan
- 26. Running this as a team workshop
- 27. What this method does not solve
- 28. Glossary
- 29. Sources and attribution

PART I

The method

How to build an argument that lands the first time.

2. Why the pyramid: the underlying argument

2.1 Readers group ideas whether you want them to or not

The human mind cannot hold many loose items at once. Faced with nine findings and no hierarchy, a reader starts inventing categories in order to retain them: by topic, by order of appearance, by whatever worries them most. That improvised grouping almost never matches yours.

Hence the central claim: if you do not impose a structure, the reader imposes a worse one. The pyramid is not a stylistic preference. It moves the cost of ordering from the reader to the writer, who is the one with the context to do it well.

2.2 People read to answer a question

Nobody reads a memo out of curiosity. They read it because they have an open question — do I approve this, should I worry, what do I do — and they are looking for the answer. A well-structured document reads like a dialogue: the idea above raises a question in the reader's mind, and the level immediately below answers it.

When that dialogue works, the reader can stop at any depth and still hold a complete, coherent version of the argument. A director reads three lines and can decide. An analyst goes three levels down and finds the detail. The same document serves both.

2.3 What a quantitative profile gains specifically

The typical failure of a technical analyst is not a lack of rigour. It is reproducing, in the document, the order in which the work was done: data, then methodology, then results, with the conclusion at the end — the order of a paper. That order is right for proving something and terrible for deciding something.

The pyramid inverts the sequence without sacrificing rigour. The conclusion leads; the rigour is still there, but organised as support for the conclusion rather than as the road towards it. In practice the difference is that in a twenty-minute committee you get to say the thing that matters.

3. Anatomy of the pyramid

A well-built pyramid satisfies three rules at once. Each is verified differently, so it helps to keep them apart. The single idea at the apex is called the governing thought, or main message: the statement that governs everything below it.

3.1 The vertical rule: every idea summarises the ones below it

An idea at any level must be the summary of the ideas immediately below. Not their title, not their category: their summary. If you deleted the level below entirely, the idea above should still communicate the message.

┆ Fails: "We identified three problems in the portfolio."

┆ Works: "The deterioration is concentrated in a single sector and is reversible through limit adjustments."

The first version forces you to read three bullets before you know anything. The second states the message; the three bullets now support it, which is their actual job.

3.2 The horizontal rule: ideas at the same level are of the same kind

Items hanging from the same idea must belong to one logical category: all reasons, or all steps, or all problems, or all criteria. Mixing a reason with a step and a recommendation produces a list that feels disordered even when every item is true.

Quick test: try to give the group a plural noun. "Three reasons", "four steps", "two constraints". If the only noun that fits is "things" or "points", the group is not homogeneous.

3.3 The order rule: every group follows an explicit logic

Within a group the order cannot be arbitrary. There are only three legitimate orders, and each one reveals the kind of analysis you did. Section 7 develops them.

4. MECE: the discipline of grouping

MECE stands for Mutually Exclusive, Collectively Exhaustive. A group of ideas is MECE when the items do not overlap and, together, they cover the whole relevant space.

4.1 The two halves, separately

- Mutually exclusive: no item contains or repeats part of another. If explaining point two means going back over point one, they overlap.
- Collectively exhaustive: nothing the reader would consider part of the problem is left out. The honest test is to ask what the best-informed sceptic in the room would say is missing.

4.2 Cuts that really are MECE

Robust groupings usually come from a pre-existing, verifiable structure rather than from brainstorming:

- An algebraic identity: expected loss decomposes into probability of default, loss given default and exposure. MECE by construction.
- A partition of the population: by segment, vintage, geography or rating band. MECE if the categories do not overlap and add to the total.
- A process sequence: origination, monitoring, collections. MECE if the process is well defined.
- A dichotomy: internal versus external, controllable versus not, short term versus long term.

4.3 The three usual failures

1. The "other" bucket. It almost always means the cut was badly chosen. If "other" is more than ten per cent, redo the cut.
2. Hidden overlap. "Data problems", "system problems" and "governance problems" look clean, but a field mis-captured by a system with no owner falls into all three.

3. False exhaustiveness. You cover everything you analysed, not everything that exists. A senior reader spots the gap in seconds and distrusts the rest from that point on.

An honest caveat: MECE is a verification standard, not a publication requirement. Forcing an artificial MECE grouping produces rigid, unnatural documents. Use it to test a grouping that already makes sense, not as a template to generate one.

5. The introduction: the SCQA pattern

SCQA stands for Situation, Complication, Question, Answer. It is the mechanism that installs, in the reader's mind, the question your document answers. Without it, an opening conclusion sounds abrupt or presumptuous.

5.1 The four elements

Situation

A fact about the subject that the reader accepts without argument. Not general background, not history: the point of agreement. It must be uncontroversial enough that the reader nods along. If they can object to the situation, you are not building a base, you are opening a front.

"The large corporate rating model was approved in 2021 and now covers sixty per cent of wholesale exposure."

Complication

What disturbed that stable situation and why we are having this conversation. It is the engine of the document: a regulatory change, an observed deterioration, a deadline, a validation finding, a pending decision.

"Extending it to Mexico requires recalibrating over a window that includes 2020, and internal validation found that the current calibration understates default in the mid segment."

Question

The logical consequence of the complication. In a well-built document it is not written: it arises on its own in the reader's mind. Writing it out is optional and sometimes useful in long documents, but if you have to force it, the complication is not the right one.

"Do we recalibrate, rebuild, or apply a temporary conservative adjustment?"

Answer

Your conclusion. It is the apex of the pyramid — the governing thought — and the first substantive paragraph of the document. It must be a statement with content, not a promise of content.

"We recommend recalibrating with a transitional conservative margin: it preserves the regulatory date, costs a fraction of a rebuild, and the impact on risk-weighted assets is bounded."

5.2 The four openings and when to use each

S-C-Q-A is the standard order, but it can be rearranged depending on what you know about the reader. The content does not change; the entry point does.

Opening	Order	When it works
Standard	Situation → Complication → Answer	A neutral audience that needs context. The default case.
Direct	Answer → Situation → Complication	The reader knows the problem and is short of time. Crowded committee agendas.
Concerned	Complication → Situation → Answer	There is real urgency and you want it felt. Deterioration, incident, date at risk.
Sceptical	Question → Situation → Complication → Answer	The reader doubts the framing itself. Naming the question disarms resistance.

5.3 Introduction failures

- The résumé opening: explaining everything you did before saying what you found. The reader does not need to know how hard you worked.
- An over-long situation: three paragraphs of background before the complication. The situation should take two or three lines.
- A missing complication: if you cannot name what changed, there was probably no need for the document.
- A deferred answer: "conclusions are presented in section 5". That is not a document, it is a table of contents with filler.

6. Building the pyramid: two routes

There are two ways to arrive at the structure, top-down and bottom-up. They are used at different moments and both are legitimate.

6.1 Top-down (when you already have the conclusion)

1. Write the subject of the document in one line.
2. Write the question the reader has about that subject.
3. Write your answer to that question, even provisionally.
4. Check the introduction: can you state a situation the reader accepts and a complication that generates exactly that question? If not, your answer does not address what they want to know.
5. Anticipate the next question your answer raises — usually "why?" or "how?" — and write the three or four points that answer it.
6. Repeat down each branch until the next question no longer interests the reader.

6.2 Bottom-up (when you have loose findings)

1. List every point you want to communicate, unedited, one per line.
2. Identify the relationships: which cause which, which are instances of the same phenomenon, which are consequences.
3. Group the ones that share a nature and give each group a plural noun: reasons, symptoms, options, constraints.
4. For each group write one sentence summarising what they say together — not what they are, what they say. That sentence moves up a level.
5. Repeat until you reach a single idea. That is your answer.
6. Now build the introduction backwards and check that the question it generates is the one your idea answers. If they do not match, the grouping broke somewhere and needs redoing.

In practice the two are used together: you come down from a hypothesis, hit evidence that does not fit, work up from the findings, and adjust the apex. What matters is closing the loop at least once before publishing.

7. Ordering ideas within a group

There are only three logical orders. Any other order is really the order in which things occurred to you. Choosing one of the three consciously is also a way of auditing your own analysis: the order reveals the kind of reasoning you did.

7.1 Time order (chronological or causal)

The items are steps that happen one after another, or linked causes. Used when describing a mechanism, a flow or an implementation sequence.

Validity test: if you swap two items, does the argument stop making sense? If you can reorder freely, it was never a process — it was a list wearing a process costume.

7.2 Structural order (parts of a whole)

The items are components of a structure: regions, business lines, terms in a formula, layers of a system. Used when decomposing something that exists.

Validity test: do the parts add up to the whole without overlap? If not, the structural cut is wrong and the flaw travels upward.

7.3 Degree order (by importance)

The items share a property and are ordered by how much of it they have: most severe to least, largest impact to smallest. Used when ranking or prioritising.

Validity test: is there a metric, even a qualitative one, that justifies the order? If you cannot name it, the order is opinion dressed as analysis. In a risk function this is easy to do well: order by exposure, by capital impact, by likelihood of materialising.

Rule of thumb: if a group admits none of the three orders, the group is malformed. Regroup before writing any more. It is cheaper to break the structure now than to discover the gap in the committee room.

8. Summarising groups: where most documents are lost

This is where most people abandon the method without noticing. They build the pyramid, group well, and then label each group with an empty tag. The result looks structured and communicates nothing.

8.1 Intellectually blank assertions

These are sentences that announce the existence of content without delivering it. You recognise them because the reader cannot do anything with them.

Blank	With content
The company has three objectives.	The three objectives are mutually incompatible at the current budget.
We analysed model performance.	The model discriminates well but calibrates badly: the ranking holds, the level does not.
There are several regulatory considerations.	Neither route meets the regulatory deadline without a transitional adjustment.
We identified improvement opportunities.	Two controls have had no formal owner since the reorganisation.

8.2 The two summary rules

- Action groups (steps, recommendations, activities): the idea above must state the result of carrying out all the steps. Not "four actions for the project", but "these four actions close the project before the regulatory date".
- Situation groups (findings, facts, evidence): the idea above must state the inference that follows from all of them. Not "three validation findings", but "all three findings point to the same origin: the calibration window".

8.3 The "So what?" test

After writing any summary idea, ask it out loud: so what? If you have an answer, that answer was the idea you should have written. Repeat until the question stops producing anything new.

"Mid-segment arrears rose 40 basis points." → So what?

"It is the third consecutive quarter of increases." → So what?

"The deterioration is structural, not seasonal, and the current cut-off no longer protects the target return." ← This is the idea.

This is the highest-return exercise in the guide and the most uncomfortable, because it exposes the times when you do not actually have a conclusion, only numbers. If you run it and arrive nowhere, the problem is not the writing: the analysis is missing.

9. Deduction and induction

A group of ideas holds together in one of two ways, and it is worth knowing which one you are using, because they summarise differently.

9.1 Deduction

A chain: a statement, a second statement about the first, and the conclusion that follows from both. Each step depends on the one before.

The calibration window does not cover a full cycle. → A model calibrated without a full cycle understates default in a downturn. → The current calibration understates default in a downturn.

It is rigorous and very convincing in detail. The cost is that the reader must hold the whole chain in mind before receiving the message, and the argument breaks entirely if one link fails.

9.2 Induction

A set of items of the same kind from which a common conclusion follows. There is no chain: each item supports the conclusion on its own.

The window does not cover a cycle. The sample excludes the mid segment. The cut-off has not been reviewed in three years. → All three share one origin: the calibration is anchored to conditions that no longer hold.

It is easier to follow, more robust — if one item falls, the conclusion survives — and it lets you state the conclusion first.

9.3 Which to use

Rule of thumb: induction at the top levels, deduction at the lower ones. Senior readers want the pattern; technical reviewers want the chain. A document that uses deduction at the apex forces everyone to read everything, which is exactly what you are trying to avoid.

10. Before structuring: defining the problem properly

A good share of badly structured documents do not have a structure problem at all: they answer, precisely, a question nobody asked. This section is the antidote.

10.1 The definition sequence

Before writing anything, answer these five in writing. It takes ten minutes and saves entire rewrites.

1. Starting point: what was the stable situation, and who is the actor?

2. Disturbing event: what changed?
3. Undesired result (R1): where are we today as a consequence?
4. Desired result (R2): where do we want to be, on what metric and by when?
5. Question: what exactly does the reader need to decide or know?

The distance between R1 and R2 is your problem. If you cannot measure that distance, the document will float. And if the desired result is "to improve", you do not have an objective, you have an aspiration.

10.2 Diagnostic frameworks

When the question is "why is this happening", decompose the problem into a tree before investigating. Three ways to build one:

- By mechanism: reproduce how the process or system actually works and locate where it breaks. The most powerful and the most demanding of business knowledge.
- By possible cause: enumerate the categories of cause that could produce the observed effect, then rule them out with evidence. Useful when the mechanism is unclear.
- By undesirable elements: list what should not be happening and trace each to its origin. Useful in incidents.

10.3 Logic trees for decisions

When the question is "what do we do", the tree decomposes options rather than causes, with explicit evaluation criteria. Its advantage over a narrative analysis is that it forces you to declare the criterion before knowing the outcome, which is the only honest way to compare alternatives.

11. Getting the pyramid onto the page

The mental structure has to become visible. A well-thought, badly laid-out document reads as poorly as a badly thought one.

11.1 Headings that are ideas, not topics (message titles)

A heading should stand alone and communicate something. "Methodology" is a topic. "The recalibration uses a window extended to twelve years" is an idea. If the reader scans only your headings, they should come away with the whole argument. That is the standard.

The uncomfortable corollary: a twenty-page document should have eight to twelve headings with real content, not forty generic labels.

11.2 Executive summary

Not an abbreviated version of the document, and not a list of sections. It is the apex plus the first level of support, in prose. Four to eight lines. If the reader reads only that, they should be able to take the decision and explain it to someone else.

A brutal test: hand only the executive summary to someone unfamiliar with the subject and ask them what decision is being requested and why. If they hesitate, rewrite it.

11.3 Transitions

Moving from one block to the next, the reader needs to know where they are. One sentence is enough: pick up the idea you just closed and announce the question the next one opens. What does not work is the decorative transition that says nothing ("The following section presents...").

11.4 Numbering and visual hierarchy

- Visual hierarchy must match logical hierarchy. Two ideas at the same level look the same; a subordinate idea looks subordinate.
- Do not go past three levels of depth in a business document. If you need a fourth, that material is an annex.
- Annexes exist to move out of the body what supports but does not argue: result tables, technical specification, robustness checks. Their existence does not excuse you from summarising their message in the body.

12. The pyramid in presentations

Adapting to slides is where the method pays off most in this context, because a committee does not read: it listens and looks.

12.1 The slide title is the message (action title)

The rule, known as the action title, is simple and almost nobody applies it: the title of each slide is that slide's conclusion, written as a full sentence. The body of the slide is the evidence for the title.

Weak title: "Arrears evolution by vintage"

Correct title: "Vintages originated after the policy change deteriorate 30% faster"

The practical consequence: if you read only the titles of your deck aloud, in order, a complete and coherent argument should come out — that is what is called the storyline. It is the best diagnostic there is for a deck. Do it before every committee; it takes ninety seconds.

12.2 One idea per slide

If the title needs a conjunction to be true — "and also" — it is two slides. The temptation to compress comes from wanting to look efficient; the effect is that the committee gets lost and asks questions that were already answered.

12.3 Charts

- The chart title states the finding, not the variable plotted.
- Mark visually what you want them to see. If everything has equal visual weight, the reader chooses what to look at.
- A chart that needs spoken explanation is useless in a document that will circulate without you. And in a bank, everything circulates without you.

13. Adapting to a banking and regulatory environment

The method was designed for consulting. Transplanting it unchanged into a supervised bank creates real friction worth anticipating.

13.1 When the reader is the regulator

A supervisor does not read to decide quickly: they read to verify compliance and traceability, and they often expect a prescribed order of exposition. You cannot freely reorder there.

The solution is not to abandon the pyramid but to double the entry: a pyramid-format summary at the front, and the body in the order the supervisor expects. The summary gives you clarity; the body gives you compliance. No supervisor has ever objected to being told the conclusion first.

13.2 When the conclusion is unwelcome

The most common objection to "conclusion first" is political: if the message is bad, saying it in line one feels blunt. The objection is legitimate and the answer is that the order does not change, the framing does — the conclusion still goes first, but stated together with the way out.

Blunt: "The model will not pass validation."

Pyramid with framing: "The model will not pass validation in its current form; a transitional conservative adjustment preserves the regulatory date while it is recalibrated."

Both say the same thing. The second is just as honest and leaves the reader holding a decision rather than a problem.

13.3 Dual reporting lines

When the same matter goes to two bosses with different interests, the temptation is to write something neutral that offends neither. The usual result is a document that serves neither.

The pyramid helps here: the apex — the fact and the recommendation — is the same for both; what changes is the introduction, because the relevant complication differs for each. One document, two entrances. That is not double-speak: it is recognising that each reader arrives with a different question.

PART II

Applied cases

Eight recurring situations, each with its structure and a worked example.

14. Three worked examples

14.1 Portfolio deterioration memo

Typical version, written in work order:

"The quarterly review of the wholesale portfolio was performed. 1,240 borrowers were reviewed with data as of June. The methodology consisted of migration rates by rating and vintage analysis. Results show variations across several segments. Conclusions and recommendations are presented in section 4."

Pyramid version:

"This quarter's deterioration is concentrated in a single sector and is reversible without changing risk appetite.

Export manufacturing accounts for 70% of downgrades, against a stable picture in the rest of the book. The deterioration originates in borrowers approved under the previous cut-off, not in the performance of the current model. A sector limit adjustment plus a review of the twelve largest exposures contains the effect without restricting origination elsewhere.

We need approval of the sector adjustment at the 15 September committee."

The order changed, not the content. The first version demands twenty minutes of reading to arrive at the same place; the second lets the reader decide in thirty seconds and go deeper if they want to.

14.2 Model redevelopment proposal

SCQA introduction:

- Situation: the current model was approved five years ago and underpins the capital requirement for the wholesale portfolio.
- Complication: annual validation found understatement in the mid segment, and the regulatory calendar does not allow a full rebuild before the committed date.
- Question (implicit): recalibrate, rebuild or adjust?
- Answer: recalibrate with a transitional conservative margin.

First level of support — three reasons, in degree order, heaviest first:

1. It is the only one of the three routes that meets the committed date.
2. The diagnosis points to calibration, not specification: a rebuild would attack a problem we do not have.
3. The impact on risk-weighted assets is bounded and quantified, and the transitional margin is withdrawn at the next full recalibration.

Note that the three reasons are of the same kind, do not overlap, and together cover the three dimensions a committee evaluates: time, diagnosis and capital impact. That is MECE applied to something that already made sense.

14.3 A short email to a functional manager

The method works just as well in five lines. A well-structured escalation email:

"I need your decision on the scope of the model extension before Friday.

The data team confirmed that mid-segment information is only available from 2019. That leaves two options: extend with the short window and document the limitation, or wait for the history to be rebuilt and move the date by a quarter.

My recommendation is the first: the limitation is documentable, the delay is not. If you prefer the second, I need to tell the project team this week."

Request first, context second, explicit recommendation, and a clear consequence on each branch. No ornament.

15. Data quality when someone else owns the data

This is probably the most frequent and worst-handled situation in an analytics function: you depend on information produced by someone else, the information arrives flawed, and you have no authority over whoever generates it. The typical outcome is an email listing defects that the owning team receives as an accusation, answers defensively, and nobody fixes anything.

15.1 Why the intuitive approach fails

The intuitive approach puts the defect at the apex: "we found 43 inconsistencies in the collateral file". That apex turns the recipient into a defendant and you into an uninvited auditor. It also contains no request: the reader does not know what is expected of them, so the email gets filed.

The correct apex is not the defect. It is the decision or deliverable at risk, and the bounded action that unblocks it. The defect moves to the second level, where it belongs: it is evidence, not the message.

15.2 Recommended structure

- Situation: name the shared process in terms the other party recognises as theirs and as legitimate. "The monthly collateral file feeds the loss-given-default calculation and the September regulatory report."
- Complication: the measured defect plus its dated consequence. Measured means a rate and a volume, not adjectives. "Twelve per cent of records arrive with no valuation date; at that level the calculation cannot be signed off before 15 September."
- Answer: the smallest request that solves the problem, addressed to one person, with a date. Not "improve the quality of the file", but "populate the valuation date on the 340 records in the annex before 5 September".
- Support: what you have already done. This is not self-promotion: it is what stops the other party feeling you are handing them your work. "We applied a temporary control that imputes the date by segment average; it works for internal monitoring, not for the regulatory calculation."

15.3 Rules that change the outcome

1. Always quantify. A defect with no rate and no impact is an opinion, and opinions do not move another team's priorities.
2. Translate the defect into their consequence, not yours. "We have to reprocess" means little to them; "the figure you sign will go out qualified" means a great deal.
3. One request per counterpart. An email with nine requests becomes zero requests actioned.
4. Prioritise visibly. If you have 43 defects, group them by origin — capture, system, definition — and say which block and which do not. Half the time, three defects explain eighty per cent of the impact.
5. Offer the path of least effort. If you can hand them the exact list of records to correct instead of asking them to find it, do it. Half an hour of your time costs less than two weeks of back and forth.
6. Leave a record without an accusatory tone. Documentation is necessary and needs no adjectives: dated facts are enough, and they age far better if the matter escalates.

15.4 The escalation path

Escalating is neither a punishment nor a betrayal; it is a normal mechanism when two functions hold different priorities. What makes it toxic is doing it by surprise. The clean sequence has three rungs and at every one the recipient knows what comes next:

1. Direct request with a date and a consequence. Nothing more.
2. A reminder with the impact updated and an explicit note that, absent a reply, the matter goes to the relevant forum. Not as a threat — as information.
3. Escalation with options. When you escalate, do not bring the problem: bring two routes with their costs. An escalation that only reports leaves your manager with work; one that proposes leaves them with a decision.

15.5 Worked example

Typical version:

"Good morning. Attached is the result of the quality review of the August collateral file. 43 issues of various types were detected and are detailed in the annex. It is worth mentioning that these issues have recurred in previous months. I look forward to your comments."

Pyramid version:

"I need the valuation date for 340 collateral items before 5 September in order to sign off the loss-given-default calculation for the regulatory close."

The August file arrived with that field empty on 12% of records. We applied a segment-level imputation that works for internal monitoring but is not admissible for the reported calculation.

The annex lists the exact 340 records, sorted by exposure. The first 60 account for 70% of the impact; if time is tight, we can proceed with those.

The remaining issues this month do not block the close; I will send them separately next week so they can be reviewed properly at source."

What changed: the apex moved from the defect to the decision at risk, the request became bounded and dated, the search work is already done, a partial route is offered if time runs short, and the urgent was separated from the important. Same content, different outcome.

16. Getting a decision in a globally centralised organisation

When the decision is not taken where you are, the problem stops being technical and becomes a competition for attention. Whoever decides at the centre covers twenty countries, and your request competes with the other nineteen. The pyramid helps because in that setting you have one paragraph before you are classified.

16.1 Changing the apex

The intuitive apex is "Mexico needs X". From the centre, that sentence reads as a local exception: cost for them, benefit for you. It is the category with the lowest approval rate there is.

The apex that works reframes the same request as the solution to a problem the centre already recognises as its own. This is not manipulation: it is choosing, among several true things about your request, the one that is relevant to the person deciding.

Local: "We need an exception to the model standard for the Mexican portfolio."
Central: "The global standard does not cover markets with short data history; Mexico can be the pilot that closes that gap before three other countries open it."

16.2 The five elements of a case to the centre

1. Anchor in a stated priority. Cite the group objective, programme or policy your proposal aligns with. If none exists, your proposal will compete without backing.
2. Translate into the centre's metric. What worries you in local currency and in operations matters to them in capital, in regulatory risk, or in consistency across entities. Translate.
3. Ask for the smallest possible decision. Do not request the new policy: request the pilot, the bounded exception with a review date, the temporary waiver. Reversible things are far easier to approve.
4. Handle the precedent. Whoever decides centrally is always thinking "if I give it to one, they will all ask". Get ahead of it: define the criterion that makes your case a category rather than a favour. "Applicable to entities with less than five years of history in the segment" is a criterion; "applicable to Mexico" is a favour.
5. State the cost of not deciding, in their currency. The argument that moves a banking group most is local regulatory risk, because the centre does answer for it. Use it when it is true, and do not inflate it when it is not: credibility lost by exaggerating once does not come back.

16.3 Who you are actually addressing

Before writing, answer three questions on paper: who signs, who can veto without signing, and who is affected if this is approved. Most proposals fail because of the second group, who were not consulted and find out in the forum.

Where there is dual reporting — a functional line to the centre and an entity line locally — tending only one of them is an expensive and very common mistake. The functional line gives you technical viability; the entity line gives you political backing. You need both, and each needs to hear its own version before the forum, never during it.

16.4 The format that survives

One page. Request and decision sought in the first paragraph. Three reasons. Quantified impact. The date after which the decision stops being useful. Annexes for everything else. If your proposal needs fifteen slides to be understood, it is not ready for the centre; it is ready for your team.

17. Building a business case for funding

A business case is a sales document in the format of a technical one. Whoever reads it has seen dozens and has developed a detector for inflated benefits. The pyramid structure helps, but here credibility matters as much as order.

17.1 The apex

The apex is the decision you are asking for and its amount. Literally: what you approve, what it costs, what you get, by when. A business case that opens with "project background" lost the reader in line one.

"I am requesting approval of a budget of X to automate the regulatory reconciliation process over twelve months. It frees the equivalent of two manual positions, closes the audit finding open since March, and pays back in under two years."

17.2 Structure

1. Today's cost, quantified. Not the problem in the abstract: what doing nothing costs, in hours, in losses, in open risk. This is the number that carries the whole document and the one most often skipped out of laziness.
2. The options considered, including doing nothing. A business case with a single option reads as a decision already taken looking for a stamp. Three options with explicit criteria read as analysis.
3. The recommended option and why it wins under those criteria.
4. Cost and benefit with visible assumptions. Every assumption behind the benefit, written and attackable. If the reader cannot see the assumptions, they assume they are inflated — and they are usually right.
5. What happens if it is not approved. A concrete consequence, not a dramatic one.
6. The request and the real deadline for the decision.

17.3 Credibility rules

- Benefit in the decision-maker's currency. Hours saved convinces a head of operations; capital or regulatory risk reduction convinces a committee. The same project is argued differently depending on who signs, without lying in either version.
- Honest ranges beat false precision. "Between 1.5 and 2.5 full-time equivalents, depending on phase two scope" is more credible and more defensible than an exact number nobody can stand behind.
- Quantify the cost, not just the benefit. Include what usually gets omitted: the time of a team that is already busy, year-two maintenance, the learning curve.
- Name the main risk yourself. If you do not put it on the table, someone will put it there in worse terms. Naming it with its mitigation gives you control of the framing.
- Answer "why now". Almost no business case is rejected for being bad; it is rejected for being postponable. If you cannot explain why this quarter and not next, expect it to slip.

18. Difficult conversation downward: asking someone on your team to change a behaviour

When one person's behaviour damages the team, the cost of not having the conversation is always paid by the rest of the team, and paid in silence. Postponing it is not kindness: it is transferring the problem to the people who did not cause it.

18.1 Preparation: three columns before you speak

- Observable fact. What happened, when, in what context. Verifiable, specific, adjective-free: something a third party who was there would describe the same way.
- Impact. What consequence it had on the work or the team. Not on your feelings: on the work. "Two people stopped bringing items to the session" is impact; "it makes me uncomfortable" is not.
- Expected change. What specifically must be different and what success would look like in four weeks. If you cannot describe it, you are not ready for the conversation.

If filling in the three columns reveals that you have impressions and no facts, the problem may still be real — but you need to observe first. A conversation built on perceptions turns into an argument about whether the perception is fair, and that goes nowhere.

18.2 Behaviour, not identity

This distinction decides the outcome. A criticism of behaviour admits correction; a criticism of identity admits only defence, because you are telling someone who they are.

Identity (closes down)	Behaviour (opens up)
You are arrogant with the team.	In the last three sessions you interrupted the analysts before they finished presenting.
You are not collaborative.	You closed the last two shared deliverables without telling the two people who depended on them.

Identity (closes down)	Behaviour (opens up)
You have a bad attitude.	When you were asked to adjust the report, you replied in the team chat that it was a waste of time.

18.3 The script, in five beats

1. A direct opening, no circling and no filler praise. "I want to talk about something that is affecting how the team works." Praise beforehand makes the person brace for the blow and devalues any future praise from you.
2. The fact. One or two, the clearest. Not the full list: a long list feels like a case file and pushes the conversation towards defence.
3. The impact. Concrete and in terms of the work.
4. Real listening. Questions and silence. This is where coaching practice meets this method: ask, and hold the pause. Information may surface that changes your reading, and if it does, it has to be genuinely incorporated, not performed.
5. A verifiable agreement and a follow-up. What changes, how we will notice, when we review it. Without a review date, the conversation did not happen.

18.4 The three reactions and what to do

- They deny the fact. Return to the fact, not the judgement. If the disagreement is about what happened, that is the thing to resolve before anything else.
- They fall apart. Slow down, hold the message. Withdrawing it out of discomfort is the most common mistake and guarantees having the conversation again in worse conditions. You can be firm and considerate at once.
- They counter-attack with their own grievances. Acknowledge whatever is legitimate, schedule it separately with a date, and return. Merging the two topics leaves both unresolved.

18.5 Documentation

A short email after the conversation — what was agreed, the review date, nothing more — protects both parties and prevents a later argument about what was said. It is not a case file and needs no formal tone; it is shared memory. If the situation escalates later, that record is the difference between an orderly process and one person's word against another's.

19. Difficult conversation upward: when the problematic behaviour is your manager's

This is a different case and should not be treated as the previous one with the roles reversed. Downward you request a change with authority; upward you have no such lever and the cost of error is asymmetric. That does not mean staying quiet: it means preparing better.

19.1 What changes from the previous case

- You do not request a behaviour change: you present an effect and make a small request. The conclusion still goes first, but the conclusion is the request, not a diagnosis of their character.
- Impact is measured in results, not in mood. "When priorities change mid-week without notice, two deliverables slip" can be discussed. "The team feels demotivated" invites a challenge to your reading of the team.
- You ask permission for the conversation. One sentence changes the whole tone: "Can I raise something about how we are working? It is uncomfortable but I think it is useful for you to know."
- You offer a way out. A concrete proposal that makes it easy to do things differently, not just the observation. Without a proposal, the conversation is a well-drafted complaint.

19.2 Structure

"Do you have ten minutes? I want to propose an adjustment to how priority changes reach the team.

Over the last three weeks, two committed deliverables slipped because the focus changed mid-week. The team covered the urgent items, but we were late to finance twice.

My proposal is to agree priority changes on Mondays, and for anything arriving later to wait unless it is regulatory. If something comes in outside that window, I tell you what drops, you decide, and we are both covered.

Shall we try it for a month?"

Note what it does not contain: no adjectives about them, no generalisation, no accumulated grievances. A measured effect, a proposal and a trial period. That is hard to argue with, and if it is still refused, you have learned something valuable about the room you actually have.

19.3 Calibrating the risk

- Channel and timing: never in public, never in the heat of the moment, never in writing as first contact. Email closes options because it persists; a conversation leaves room for both parties to adjust.
- One thing per conversation. The temptation to use the moment to raise everything accumulated is what turns a reasonable request into a grievance.
- When it is not worth it: when the pattern is structural and sustained, when you have already tried this approach with no movement, or when the personal cost exceeds the expected benefit. Recognising that is not cowardice; it is calculation. What is not defensible is never trying and always complaining.
- Escalating above your manager or going to human resources carries a real and often irreversible cost to the relationship. It is the right route for conduct that crosses an ethical or dignity line, and an expensive route for management-style problems. Telling which is which, before acting, is part of the job.

20. Structuring a "no"

Saying no badly is one of the most constant sources of overload. A no wrapped in three paragraphs of context reads as "let me see", and comes back the following week converted into a commitment.

20.1 The rule

The no goes first, in the first sentence, with no preamble. Then the reason, then the alternative. That order respects the other person more than the long version: it lets them stop reading and look for another route immediately, instead of waiting for an answer that is already settled.

20.2 Three kinds of no, and the cost of confusing them

- Capacity no. You could do it, but something else drops. Never phrase it as "I can't": phrase it as an explicit trade. "Yes, if we push the X delivery by two weeks. Shall we?" This returns the decision to whoever has the authority to prioritise, which is where it belongs.
- Scope no. It does not belong to your team, or it belongs to another that will do it better. Here it matters to name who does, because a no with no destination reads as a dodge.
- Judgement no. It can be done and it should not be. This is the most uncomfortable one and the one that most defines you professionally. It requires an explicit technical reason and a willingness to hold it. It is also the only one that admits no soft alternative: if you believe something should not be done, offering a reduced version abandons your own argument.

20.3 Structure

1. The no, explicit, in the first line.
2. One reason, the real one, verifiable. Only one: three reasons read as accumulated excuses, and the other party only has to knock one down to reopen the matter.
3. The alternative, if there is one: another route, another owner, another date, a smaller scope.
4. What would make it a yes. This is what separates a collaborative no from a door slammed shut, and it often unlocks the real conversation.

20.4 Example

"We are not going to be able to include the sensitivity analysis in Thursday's version.

The team is closing the regulatory exercise that same week and there is no room without putting that date at risk.

Two options: we deliver it complete on the 12th, or Thursday goes out without sensitivity and the annex follows. My recommendation is the second, because the committee can decide without that analysis.

If sensitivity is essential for Thursday, we would need to move the regulatory exercise, and that does need your sign-off."

21. Handing over a project or a system

Almost all handovers are documented badly for the same reason: the person leaving documents what they know, not what the other person will need. And what the other person will need is, above all, what you no longer notice because it runs automatically in your head.

21.1 The apex

A handover does not end when you deliver the folder. It ends when the receiver runs a full cycle without you and it goes well. Defining it that way changes everything else: the central deliverable is not a document, it is a supported sequence with a declared cut-off date.

21.2 What the document must contain

1. What it is and why it exists, in one paragraph. Who uses it and what decision depends on it.
2. An honest current state. What works, what is half-built, and what is broken and running on a patch. Hiding technical debt does not remove it: it only guarantees the receiver finds it at the worst moment with your name on it.
3. How it is operated. Routines, calendar, inputs, dependencies, who delivers what and when. With names and specific days.
4. Who is who. Not just contacts: what each person cares about and where there is history. This is what never gets written and what is needed most.
5. Tacit assumptions. The three to five "this is always done this way because..." that live only in your head. It is the most valuable section of the document and the one that takes most effort to recover.
6. Known risks and open decisions, with your recommendation on each.
7. What to do when it fails. The three or four frequent failure modes and their fix, written as executable instructions.

21.3 The transfer sequence

Four stages, each with at least one real cycle. None is skipped even if the receiver is experienced, because what is being transferred is not technical capability but context.

1. You do it, they watch.
2. They do it, you watch and do not intervene except on material error.
3. They do it, you are available but not present.
4. They do it. A declared cut-off date and a short warranty period — two or three weeks — in which you answer questions but do not execute.

The cut-off date has to exist and has to be public. Without it, handovers become permanent: the receiver keeps asking, you keep solving, and six months later the system is still yours in practice and theirs on the org chart.

21.4 A living document, not a monument

An eighty-page manual does not get read. What gets used is a short operating document — five to ten pages — plus the technical files wherever they live. If you are short of time, write the tacit assumptions and the frequent failures first: they are the sections nobody else can write.

22. Explaining a technical subject in plain terms

The underlying misunderstanding is believing that simplifying means removing rigour. It does not. Simplifying well is changing the level of abstraction without saying anything false, and it is harder than explaining in detail because it requires knowing which part of the detail that audience can do without.

22.1 The three layers

1. What this decides. The practical consequence for the person listening. It comes first and is often all they need.
2. How it works, in one sentence. The mechanism at its coarsest level, with an analogy if it helps.
3. The detail, available but not imposed. Offer it, do not unload it. "If you like, I can show you how it is calculated" beats calculating it out loud unprompted.

22.2 Rules

- Start with the consequence. "This change means we approve fewer clients in the mid segment" before anything about methodology.
- One analogy, and say where it breaks. Every analogy is false somewhere; pointing that out yourself preserves rigour and stops someone building a wrong conclusion on it.
- Replace jargon with the effect, not with a synonym. Instead of explaining what calibration is, say: "the model ranks clients from best to worst correctly, but the absolute level of risk it assigns is too low".
- Anchor numbers to something familiar. "Four in every thousand loans" communicates more than "0.4%", and "the equivalent of two team-weeks" more than "80 hours".
- Never say "it's complicated". It reads as condescension and as an excuse. If it genuinely does not fit the time available, say which part does.
- No defensive detail. A lot of long technical explanation is not aimed at being understood; it is aimed at demonstrating that you know. It shows, and it costs credibility rather than building it.

22.3 Example

Technical version, aimed at the wrong audience:

"The recalibration adjusts the intercept of the link function to align the long-run average probability of default with the observed through-the-cycle rate, without modifying the coefficients or the discriminatory power of the model."

Same truth, for a business head:

"The model still ranks clients just as well: the ones it flags as good are good. What was wrong is the level – it said 2% default where reality has been 3%. We are raising it to 3%.

For you that means some clients who pass the cut-off today will stop passing it, mainly in the mid segment. I have the estimate of how many.

What does not change is the ranking, so relative commercial decisions – who gets a larger limit – stay the same."

22.4 The test

Ask the person to explain it back to you, casually: "how would you tell your team?" What comes back is exactly what they understood. It is uncomfortable the first time and it is the only reliable diagnostic there is; everything else — nodding, not asking — is noise.

PART III

Practice

How to install the habit, in yourself and in a team.

23. Common failures in a quantitative profile

1. Reproducing the order of the analysis instead of the order of the decision.
2. Confusing precision with relevance: putting the confidence interval in the executive summary and leaving out the recommendation.
3. Topic headings instead of message headings.
4. Lists of findings with no summary idea: the structure exists but communicates nothing.
5. Long deductive chains at the top level, forcing the reader through everything before understanding anything.
6. Grouping by data convenience — whatever came out of the same query — instead of by the logic of the argument.
7. Annexes carrying the conclusion: the key message hidden in table 7 of appendix B.
8. Avoiding the conclusion out of political caution, and delivering a document that says nothing verifiable.
9. Assuming the reader knows the jargon. Every acronym is spelled out the first time, including the obvious ones in your own function.
10. Fixing the prose when the problem is the thinking. If "so what?" produces no answer, no rewrite will fix it.

24. Checklist

Use it before sending any document or deck. Ten questions, five minutes.

1. Is the first substantive sentence my conclusion, rather than an announcement that one is coming?
2. Can I state the situation and the complication in two lines each?
3. Is the question my introduction generates the one my document answers?
4. Does every group have between two and five items of the same kind?
5. Does every summary idea survive the "so what?" test?
6. Does every group follow a time, structural or degree order, and can I name which?
7. Are the groups mutually exclusive, and what would the best-informed sceptic say is missing?
8. Reading only the headings, does the full argument come through?
9. Is the executive summary enough to decide without reading the rest?
10. Is every acronym spelled out the first time it appears?

25. A 30-day practice plan

Pyramid structure is not installed by reading. These four weeks are designed around material you have to produce anyway, with no meaningful extra time.

Week	Exercise	Time
1	Take three documents you have already sent. For each, write the SCQA introduction it should have had and the conclusion in one sentence. Do not rewrite them in full.	20 min each
2	Every email longer than five lines: request or conclusion in the first line, no exceptions. Before sending, apply "so what?" to that first line.	0 extra min
3	Take your last presentation. Rewrite every title as a full sentence with a conclusion. Read them aloud in order and find where the argument breaks.	1 hour
4	Write a new document entirely with the method, top-down. Before sending it, give a colleague only the executive summary and ask what decision is being requested.	2 hours

From the second month, the only habit to sustain is the "so what?" test applied to every heading and every slide title. Everything else becomes automatic.

26. Running this as a team workshop

Handing out the document does not teach the method. What works is a short session in which each person rewrites something of their own and gets criticism on it. This ninety-minute agenda is built on one principle: eighty per cent of the time is them writing, not you talking.

26.1 Before the session

- Ask each participant to bring a document of their own that has already been sent: a memo, a long email, a summary slide. Theirs, real, not an invented example. Without that, the session does not work.
- Send the guide in advance but do not assume anyone read it. Design the session to work without prior reading.
- Eight people maximum. Beyond that, nobody gets individual criticism and the exercise dilutes.

26.2 Agenda

Block	Content	Minutes
Opening	A real, anonymised example of a badly structured document, and the question to the group: what is being asked here? Let them struggle. That bewilderment is the argument for the whole workshop.	10
Method	Only three things: conclusion first, SCQA, and the "so what?" test. Do not teach MECE or deduction and induction in the first session.	20
Exercise 1	Each person writes the conclusion of their document in one sentence, and applies "so what?" until it produces nothing new.	15
Criticism	Three volunteers read their sentence. The group asks "so what?" out loud. It is uncomfortable and it is where the learning happens.	20
Exercise 2	Each person rewrites the introduction of their document using SCQA, in four lines.	15
Close	The concrete commitment: for the next month, every email longer than five lines carries the request in the first line. Nothing else.	10

26.3 Rules for the criticism

1. The document is criticised, never the person. It is the same behaviour-versus-identity distinction from section 18, applied to a text.
2. Whoever criticises proposes a concrete alternative. "It is not clear" is not useful criticism; "I would put this sentence first" is.
3. The author does not defend during the criticism. Take notes and decide later. Defending in the moment turns the exercise into a debate and consumes everyone's time.
4. You go first, with a document of your own. If whoever runs the workshop does not face the same scrutiny, nobody else really will.

26.4 What to expect

The first session usually reveals something more uncomfortable than a writing problem: that several documents have no conclusion because none was ever reached. That is not a failure of the workshop, it is its most valuable outcome. Say so out loud at the start so it is not experienced as exposure.

The habit does not install itself in one session. What sustains the change is what you do afterwards: if you consistently return documents with the same question — what is the conclusion, and does it come first? — in two months the team writes differently. If you do not, the workshop was a pleasant afternoon with no effect.

27. What this method does not solve

It is worth being explicit about the limits, because the method is sometimes sold as a master key to communication and it is not.

- It does not improve spoken presence. Structuring an argument and holding it under interruption, pressure or hostile questions are different skills. Those are trained by repetition in front of an audience with feedback, not by writing.
- It does not substitute for analysis. If "so what?" produces no answer, what is missing is analytical work, not structure.
- It does not fix organisational politics. An impeccable document addressed to the wrong person is still a useless document.
- It does not generate business intuition. It helps you express the judgement you already have; it does not create it.

With those four caveats in mind, it remains, in my view, the highest available return per hour invested for someone technically solid whose bottleneck is that the content does not reach the decision-maker in order.

28. Glossary

These are the terms as they are used in the method and in day-to-day work in global organisations, with a short definition of what each one names.

Acronyms

Acronym	Stands for	What it is
SCQA	Situation, Complication, Question, Answer	The introduction pattern.
MECE	Mutually Exclusive, Collectively Exhaustive	Quality standard for a grouping: no overlaps, no gaps.
BLUF	Bottom Line Up Front	Writing convention: the conclusion in the first line.

Acronym	Stands for	What it is
SCR	Situation, Complication, Resolution	A three-part variant of the same introduction pattern.

Terms of the method

Term	What it names
Governing thought / main message	The single statement that governs the whole document.
Summary statement	The sentence that summarises a group of ideas and moves up a level.
Vertical relationship	Each idea answers the question raised by the one above.
Horizontal relationship	Ideas at the same level are of the same logical kind.
Top-down	Building from the conclusion, when you already have it.
Bottom-up	Building by grouping loose findings until you reach the conclusion.
Time order	The items are steps or linked causes.
Structural order	The items are parts of a whole.
Degree order / ranking order	The items are ordered by importance or magnitude.
Intellectually blank assertion	A sentence that announces content without delivering it.
So what? test	Asking until the idea carries actionable content.
Disturbing event	What broke the stable situation and motivates the document.
Undesired result (R1) / desired result (R2)	The two ends that define the problem.
Diagnostic framework	A decomposition used to find the cause.
Logic tree	A decomposition of options with explicit criteria.
Message title	A heading that states the conclusion, not the topic.
Action title	A slide title written as a sentence with a conclusion.
Storyline	The argument that reads straight through from the titles alone.
Executive summary	The apex plus the first level of support, in prose.

Terms of the working environment

Term	What it names
Data owner	Whoever is formally accountable for the quality of a source.
Escalation path	The agreed sequence for raising an issue without burning the relationship.
Business case	The document used to request budget.
Base case / do-nothing option	What happens if nothing is approved. The real comparison.
Opportunity cost	What stops being done if this is approved.
Stage gates	Splitting the investment into tranches with a review between them.
Pre-wire / socialise	Talking the matter through with each decision-maker before the committee.
Policy exception / waiver	Formal permission to depart from a standard.
Intent vs impact	Distinguishing what someone meant from what they produced.
Feedback sandwich	Praise, criticism, praise. It does not work.
Handover / knowledge transfer	Transferring a system or project to another person.
Shadowing	The stage in which the receiver observes before executing.
Runbook	The document with routine and incident procedures.
Technical debt	Temporary fixes that are still in place and accumulating cost.
Curse of knowledge	Taking as obvious what took you years to understand.
Order of magnitude	The scale of the number, before the exact number.

29. Sources and attribution

This document is original material: the text, the examples and the cases were written for this guide and do not reproduce content from any of the works cited. What it does take from them are the methods, which belong to their authors and are credited here.

The core method

The pyramid structure, the SCQA pattern, the MECE standard, the distinction between time, structural and degree order, and the problem-definition framework all come from *The Pyramid Principle: Logic in Writing and Thinking*, by Barbara Minto (Pearson / Financial Times). Anyone who wants the original formulation and its examples should go to the book; this guide does not replace it.

Other sources

Idea	Origin
The behaviour-versus-identity and intent-versus-impact distinctions, in sections 18 and 19	Difficult Conversations: How to Discuss What Matters Most, by Douglas Stone, Bruce Patton and Sheila Heen (Harvard Negotiation Project).
Diagnosis as the first component of a decision, in section 17	Good Strategy Bad Strategy, by Richard Rumelt.
The curse of knowledge, in section 22	A behavioural economics concept coined by Camerer, Loewenstein and Weber, popularised by Chip and Dan Heath in Made to Stick.
The habit of asking before solving, in the approach to conversations	The Coaching Habit, by Michael Bungay Stanier.

About the examples

All examples are invented. The figures, portfolios, models, borrowers and situations described do not correspond to any specific institution, and any resemblance to a real case follows from the fact that these problems recur everywhere.

This guide is shared non-commercially for team use. If you find it useful, the best gesture towards the authors cited is to buy their books.